

```
import warnings
warnings.filterwarnings("ignore")
```

## ▼ Importing Libraries

## ▼ Data Processing and Visualization

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

## ▼ Machine Learning Libraries

## ▼ Data Splitting

```
from sklearn.model_selection import train_test_split
```

## ▼ Data Pre-processing

```
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler
```

## ▼ Sampling

```
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import SMOTE
```

## ▼ ML Algorithms

```
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
```

## ▼ Neural Networks

```
from keras.models import Sequential
from keras.layers import Dense, Dropout
from keras.optimizers import Adam
from keras import metrics
from keras.initializers import Constant
```

## ▼ Metrics

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

## ▼ Storing Model

```
import pickle
```

## ▼ Loading Train and Test Data

```
train_data = pd.read_csv('data/Training Data.csv')
test_data = pd.read_csv('data/Test Data.csv')
```

```
train_data
```

	Id	Income	Age	Experience	Married/Single	House_Ownership	Car_Owner
<b>0</b>	1	1303834	23	3	single	rented	
<b>1</b>	2	7574516	40	10	single	rented	
<b>2</b>	3	3991815	66	4	married	rented	
<b>3</b>	4	6256451	41	2	single	rented	
<b>4</b>	5	5768871	47	11	single	rented	
...	...	...	...	...	...	...	...
<b>251995</b>	251996	8154883	43	13	single	rented	
<b>251996</b>	251997	2843572	26	10	single	rented	
<b>251997</b>	251998	4522448	46	7	single	rented	
<b>251998</b>	251999	6507128	45	0	single	rented	
<b>251999</b>	252000	9070230	70	17	single	rented	

252000 rows × 13 columns



```
test_data
```

	ID	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownership
0	1	7393090	59	19	single	rented	
1	2	1215004	25	5	single	rented	
2	3	8901342	50	12	single	rented	
3	4	1944421	49	9	married	rented	}
4	5	13429	25	18	single	rented	}
...	...	...	...	...	...	...	
27995	27996	9955481	57	13	single	rented	
27996	27997	2917765	47	9	single	rented	
27997	27998	8082415	24	5	single	rented	
27998	27999	9474180	51	13	single	rented	}
27999	28000	9250350	42	9	single	rented	

28000 rows × 12 columns



## ▼ Exploratory Data Analysis

### ▼ Checking Data Type

```
train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 252000 entries, 0 to 251999
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Id                   252000 non-null int64
1   Income               252000 non-null int64
2   Age                  252000 non-null int64
3   Experience            252000 non-null int64
4   Married/Single       252000 non-null object
5   House_Ownership      252000 non-null object
6   Car_Ownership        252000 non-null object
7   Profession            252000 non-null object
8   CITY                  252000 non-null object
```

```

9    STATE                252000 non-null  object
10   CURRENT_JOB_YRS      252000 non-null  int64
11   CURRENT_HOUSE_YRS    252000 non-null  int64
12   Risk_Flag            252000 non-null  int64
dtypes: int64(7), object(6)
memory usage: 25.0+ MB

```

We conclude that the dataset has no null values

## ▼ Data Visualization

```

def plot_hist(data1, data2, feature, labels):
    plt.hist([data1[feature], data2[feature]], bins=10, label=labels, density=True, stacked=True)
    plt.title("Defaults based on {}".format(feature))
    plt.legend()
    plt.xlabel(feature)
    plt.show()

```

```

def plot_pie(data, labels, title, axis=None, ax_index=None):
    if axis is not None and ax_index is not None:
        axis[ax_index % 13, ax_index // 13].pie(data, labels=labels, autopct='%1.1f%%')
        axis[ax_index % 13, ax_index // 13].set_title("Defaults based on {}".format(title))
    else:
        plt.pie(data, labels=labels, autopct='%1.1f%%')
        plt.title("Defaults based on {}".format(title))
        plt.show()

```

```

def get_defaults(data, feature, value, text=None):
    all_value = data[data[feature] == value]['Risk_Flag'].value_counts()
    print("Proportion of {} who default: {}".format(text or value, all_value[1]/(all_value[0]+all_value[1])))
    return all_value

```

```

risky_1 = train_data[train_data['Risk_Flag'] == 1].copy()
risky_0 = train_data[train_data['Risk_Flag'] == 0].copy()

```

```

flags = ['Risky', 'Safe']


```

## ▼ Influence on Income on Risk factor

```

plot_hist(risky_1, risky_0, 'Income', flags)

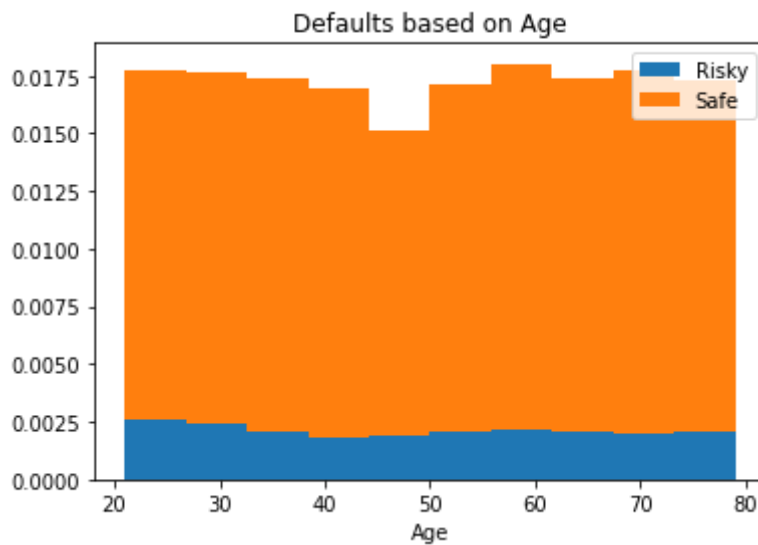
```



### ▼ Influence of Age on Risk Factor

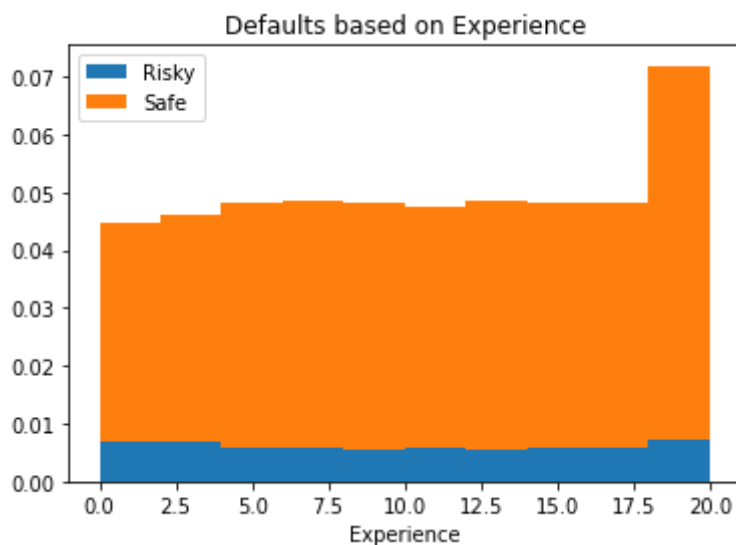


```
plot_hist(risky_1, risky_0, 'Age', flags)
```




### ▼ Influence of Experience on Risk Factor

```
plot_hist(risky_1, risky_0, 'Experience', flags)
```




### ▼ Influence of Job Years on Risk Factor

```
plot_hist(risky_1, risky_0, 'CURRENT_JOB_YRS', flags)
```



### ▼ Influence of House Years on Risk Factor

```
plot_hist(risky_1, risky_0, 'CURRENT_HOUSE_YRS', flags)
```



### ▼ Influence of Gender on Risk Factor

```
train_data['Married/Single'].value_counts()
```

```
single      226272
married     25728
Name: Married/Single, dtype: int64
```


```
all_singles = get_defaults(train_data, 'Married/Single', 'single')
```

```
Proportion of single who default: 0.1253358789421581
```

```
plot_nie(all_singles, flags[...-1], 'Singles Spread')
```

```
plot_pie(all_singles, flags[::1], 'Singles Spread')
```

Defaults based on Singles Spread




```
all_married = get_defaults(train_data, 'Married/Single', 'married')
```

Proportion of married who default: 0.10245646766169154

```
plot_pie(all_married, flags[::1], 'Married Spread')
```

Defaults based on Married Spread



## ▼ Influence of House Ownership on Risk Factor

```
train_data['House_Ownership'].value_counts()
```


```
rented      231898
owned       12918
norent_noown  7184
Name: House_Ownership, dtype: int64
```

```
all_rented = get_defaults(train_data, 'House_Ownership', 'rented')
```

Proportion of rented who default: 0.1255767621971729

```
plot_pie(all_rented, flags[::1], 'Rented Spread')
```

## Defaults based on Rented Spread




```
all_owed = get_defaults(train_data, 'House_Ownership', 'owned')
```

Proportion of owned who default: 0.08979718222635083

```
plot_pie(all_owed, flags[::-1], 'House Owned Spread')
```

## Defaults based on House Owned Spread



```
all_noown = get_defaults(train_data, 'House_Ownership', 'norent_noown')
```

Proportion of norent\_noown who default: 0.09952672605790645

```
plot_pie(all_noown, flags[::-1], 'No Rent No Own Spread')
```




## ▼ Influence of Car Ownership on Risk Factor

```
all_car = get_defaults(train_data, 'Car_Ownership', 'yes', 'Car Owners')
```

Proportion of Car Owners who default: 0.11098684210526316

```
plot_pie(all_car, flags[::-1], 'Car Owners Spread')
```

Defaults based on Car Owners Spread




```
no_car = get_defaults(train_data, 'Car_Ownership', 'no', 'Non Car Owners')
```

Proportion of Non Car Owners who default: 0.1281875

```
plot_pie(no_car, flags[::-1], 'Non Car Owners Spread')
```

Defaults based on Non Car Owners Spread



## ▼ Influence of Profession on Risk Factor

```
train_data['Profession'].value_counts()
```

Physician	5957
Statistician	5806
Web_designer	5397

Psychologist	5390
Computer_hardware_engineer	5372
Drafter	5359
Magistrate	5357
Fashion_Designer	5304
Air_traffic_controller	5281
Comedian	5259
Industrial_Engineer	5250
Mechanical_engineer	5217
Chemical_engineer	5205
Technical_writer	5195
Hotel_Manager	5178
Financial_Analyst	5167
Graphic_Designer	5166
Flight_attendant	5128
Biomedical_Engineer	5127
Secretary	5061
Software_Developer	5053
Petroleum_Engineer	5041
Police_officer	5035
Computer_operator	4990
Politician	4944
Microbiologist	4881
Technician	4864
Artist	4861
Lawyer	4818
Consultant	4808
Dentist	4782
Scientist	4781
Surgeon	4772
Aviator	4758
Technology_specialist	4737
Design_Engineer	4729
Surveyor	4714
Geologist	4672
Analyst	4668
Army_officer	4661
Architect	4657
Chef	4635
Librarian	4628
Civil_engineer	4616
Designer	4598
Economist	4573
Firefighter	4507
Chartered_Accountant	4493
Civil_servant	4413
Official	4087
Engineer	4048

Name: Profession, dtype: int64


```

professions = list(train_data['Profession'].value_counts().index)
fig, ax = plt.subplots(13, 4, figsize=(25, 52))
ax_index = 0
for profession in professions:
    all_data = get_defaults(train_data, 'Profession', profession)
    plot_pie(all_data, flags[::-1], str(profession), ax, ax_index)
    ax_index += 1
_ = ax[-1, -1].axis('off')


```

Proportion of Physician who default: 0.11918751049185831  
 Proportion of Statistician who default: 0.11557009989665863  
 Proportion of Web\_designer who default: 0.10913470446544377  
 Proportion of Psychologist who default: 0.12189239332096476  
 Proportion of Computer\_hardware\_engineer who default: 0.12844378257632166  
 Proportion of Drafter who default: 0.1128941966784848  
 Proportion of Magistrate who default: 0.12002986746313235  
 Proportion of Fashion\_Designer who default: 0.11538461538461539  
 Proportion of Air\_traffic\_controller who default: 0.1353910244271918  
 Proportion of Comedian who default: 0.11960448754516068  
 Proportion of Industrial\_Engineer who default: 0.09866666666666667  
 Proportion of Mechanical\_engineer who default: 0.11155836687751582  
 Proportion of Chemical\_engineer who default: 0.11162343900096061  
 Proportion of Technical\_writer who default: 0.134167468719923  
 Proportion of Hotel\_Manager who default: 0.13538045577443028  
 Proportion of Financial\_Analyst who default: 0.10315463518482679  
 Proportion of Graphic\_Designer who default: 0.11536972512582269  
 Proportion of Flight\_attendant who default: 0.12363494539781592  
 Proportion of Biomedical\_Engineer who default: 0.12755997659449972  
 Proportion of Secretary who default: 0.13040901007705988  
 Proportion of Software\_Developer who default: 0.1484266772214526  
 Proportion of Petroleum\_Engineer who default: 0.08510216226939099  
 Proportion of Police\_officer who default: 0.16405163853028798  
 Proportion of Computer\_operator who default: 0.12404809619238477  
 Proportion of Politician who default: 0.11225728155339806  
 Proportion of Microbiologist who default: 0.12435976234378202  
 Proportion of Technician who default: 0.12828947368421054  
 Proportion of Artist who default: 0.1226085167660975  
 Proportion of Lawyer who default: 0.1295143212951432  
 Proportion of Consultant who default: 0.1252079866888519  
 Proportion of Dentist who default: 0.109577582601422  
 Proportion of Scientist who default: 0.14432127170048106  
 Proportion of Surgeon who default: 0.11546521374685666  
 Proportion of Aviator who default: 0.13493064312736444  
 Proportion of Technology\_specialist who default: 0.08148617268313278  
 Proportion of Design\_Engineer who default: 0.1069993656164094  
 Proportion of Surveyor who default: 0.15146372507424694  
 Proportion of Geologist who default: 0.144263698630137  
 Proportion of Analyst who default: 0.12146529562982006  
 Proportion of Army\_officer who default: 0.15211328041192876  
 Proportion of Architect who default: 0.13120034356882113  
 Proportion of Chef who default: 0.12146709816612729  
 Proportion of Librarian who default: 0.11257562662057044  
 Proportion of Civil\_engineer who default: 0.1358318890814558  
 Proportion of Designer who default: 0.10917790343627665  
 Proportion of Economist who default: 0.09927837305926088  
 Proportion of Firefighter who default: 0.13578877301974707  
 Proportion of Chartered\_Accountant who default: 0.15357222345871355  
 Proportion of Civil\_servant who default: 0.11579424427826875  
 Proportion of Official who default: 0.1357964276975777  
 Proportion of Engineer who default: 0.11808300395256917


Defaults based on Physician




Defaults based on Technical\_writer




Defaults based on Technician




Defaults based on Army\_officer




Defaults based on Statistician




Defaults based on Hotel\_Manager



Defaults based on Artist




Defaults based on Architect







Defaults based on Chemical\_engineer



Defaults based on Microbiologist

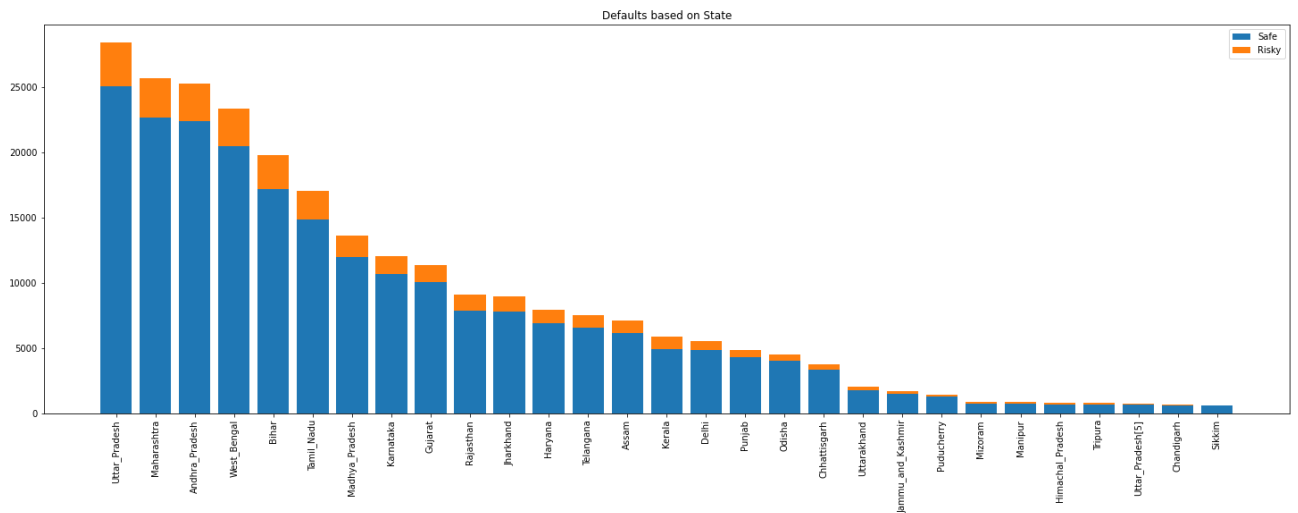


Defaults based on Analyst



## ▼ Influence of State on Risk Factor

```
states = list(train_data['STATE'].value_counts().index)
fig, ax = plt.subplots(figsize=(25, 8))
ax.bar(states, risky_0['STATE'].value_counts())
ax.bar(states, risky_1['STATE'].value_counts(), bottom=risky_0['STATE'].value_counts())
plt.title("Defaults based on State")
plt.xticks(rotation=90)
_ = plt.legend(flags[:-1])
```



## ▼ Data Preprocessing

### ▼ Remove unwanted columns

```
train_data.drop('Id', axis=1, inplace=True)
```

### ▼ Encoding Categorical Columns

```
categorical_cols = [col for col in train_data.columns if train_data[col].dtype=="O"]

label_encoder = LabelEncoder()

for col in categorical_cols:
    train_data[col] = label_encoder.fit_transform(train_data[col])
```

```
train_data
```

	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownership	Pr
0	1303834	23	3	1	2	0	
1	7574516	40	10	1	2	0	
2	3991815	66	4	0	2	0	
3	6256451	41	2	1	2	1	
4	5768871	47	11	1	2	0	

```
train_data = train_data.drop_duplicates(subset=list(train_data.columns)[: -1])
train_data.shape
```

```
(42007, 12)
```

## ▼ Splitting Dataset into Training and Validation Set

```
X = train_data.drop('Risk_Flag', axis=1)
y = train_data['Risk_Flag']
```

```
smote = SMOTE(random_state=42)
X_over, y_over = smote.fit_resample(X.values, y)
X_train_over, X_val_over, y_train_over, y_val_over = train_test_split(X_over, y_over, test
```

```
undersampler = RandomUnderSampler(sampling_strategy='majority')
X_under, y_under = undersampler.fit_resample(X.values, y)
X_train_under, X_val_under, y_train_under, y_val_under = train_test_split(X_under, y_under
```

```
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=101, s
```

## ▼ Scaling Data Values in Training Set

```
scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
X_val = scaler.transform(X_val)

X_train_over = scaler.fit_transform(X_train_over)
X_val_over = scaler.transform(X_val_over)

X_train_under = scaler.fit_transform(X_train_under)
X_val_under = scaler.transform(X_val_under)
```

## ▼ Training Model

```
def model_apply(model, X_training, y_training, X_validation, y_validation, filename):
    model.fit(X_training, y_training)
    y_pred = model.predict(X_validation)

    print("Accuracy: {}".format(accuracy_score(y_validation, y_pred)))

    print("Classification Report: \n{}".format(classification_report(y_validation, y_pred)))

    conf_matrix = confusion_matrix(y_validation, y_pred)
    print("Confusion Matrix: ")
    sns.heatmap(conf_matrix, annot=True, fmt='d')

    filename = "model/{}.pkl".format(filename)
    pickle.dump(model, open(filename, 'wb'))

logistic_regression = LogisticRegression(class_weight='balanced', penalty="l1", solver="lbfgs")
decision_tree = DecisionTreeClassifier(class_weight='balanced', criterion='gini', random_state=101)
random_forest = RandomForestClassifier(class_weight='balanced', random_state=101, n_estimators=100)
neighbours = KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2)
```

## ▼ Logistic Regression

```
print("On entire Data")
model_apply(logistic_regression, X_train, y_train, X_val, y_val, 'LogisticRegression')
```


On entire Data

Accuracy: 0.4900023803856225

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.48	0.60	6745
1	0.20	0.52	0.29	1657
accuracy			0.49	8402
macro avg	0.50	0.50	0.44	8402
weighted avg	0.68	0.49	0.54	8402

Confusion Matrix:



```
print("On Random Under Sample")
```



```
model_apply(logistic_regression, X_train_under, y_train_under, X_val_under, y_val_under, '
```


On Random Under Sample

Accuracy: 0.4939649969824985

Classification Report:

	precision	recall	f1-score	support
0	0.50	0.40	0.44	1677
1	0.49	0.59	0.54	1637
accuracy			0.49	3314
macro avg	0.49	0.50	0.49	3314
weighted avg	0.49	0.49	0.49	3314

Confusion Matrix:



```
print("On SMOTE Sample")
```

```
model_apply(logistic_regression, X_train_over, y_train_over, X_val_over, y_val_over, 'Logi
```

On SMOTE Sample  
Accuracy: 0.6044184150048187


## ▼ Decision Tree

```
print("On entire Data")
model_apply(decision_tree, X_train, y_train, X_val, y_val, 'DecisionTree')
```

On entire Data  
Accuracy: 0.6854320399904784  
Classification Report:

	precision	recall	f1-score	support
0	0.80	0.81	0.80	6745
1	0.19	0.19	0.19	1657
accuracy			0.69	8402
macro avg	0.50	0.50	0.50	8402
weighted avg	0.68	0.69	0.68	8402

Confusion Matrix:



```
print("On Random Under Sample")
model_apply(decision_tree, X_train_under, y_train_under, X_val_under, y_val_under, 'Decisi
```

On Random Under Sample

Accuracy: 0.4927579963789982

Classification Report:

	precision	recall	f1-score	support
0	0.50	0.49	0.49	1677
1	0.49	0.50	0.49	1637
accuracy			0.49	3314
macro avg	0.49	0.49	0.49	3314
weighted avg	0.49	0.49	0.49	3314

```
print("On SMOTE Sample")
```

```
model_apply(decision_tree, X_train_over, y_train_over, X_val_over, y_val_over, 'DecisionTr
```


On SMOTE Sample

Accuracy: 0.7175476314033657

Classification Report:

	precision	recall	f1-score	support
0	0.72	0.71	0.71	6761
1	0.71	0.73	0.72	6728
accuracy			0.72	13489
macro avg	0.72	0.72	0.72	13489
weighted avg	0.72	0.72	0.72	13489

Confusion Matrix:



## ▼ Random Forest

```
print("On entire Data")
```

```
model_apply(random_forest, X_train, y_train, X_val, y_val, 'RandomForest')
```

On entire Data

Accuracy: 0.8026660318971673

Classification Report:

	precision	recall	f1-score	support
0	0.80	1.00	0.89	6745
1	0.00	0.00	0.00	1657
accuracy			0.80	8402
macro avg	0.40	0.50	0.45	8402
weighted avg	0.64	0.80	0.71	8402

Confusion Matrix:



```
print("On Random Under Sample")
```

```
model_apply(random_forest, X_train_under, y_train_under, X_val_under, y_val_under, 'Random
```


On Random Under Sample

Accuracy: 0.4945684972842486

Classification Report:

	precision	recall	f1-score	support
0	0.50	0.50	0.50	1677
1	0.49	0.49	0.49	1637
accuracy			0.49	3314
macro avg	0.49	0.49	0.49	3314
weighted avg	0.49	0.49	0.49	3314

Confusion Matrix:



```
print("On SMOTE Sample")
```

```
model_apply(random_forest, X_train_over, y_train_over, X_val_over, y_val_over, 'RandomFore
```


On SMOTE Sample

Accuracy: 0.7665505226480837

Classification Report:

	precision	recall	f1-score	support
0	0.75	0.81	0.78	6761
1	0.79	0.73	0.76	6728
accuracy			0.77	13489
macro avg	0.77	0.77	0.77	13489
weighted avg	0.77	0.77	0.77	13489

Confusion Matrix:



## ▼ Nearest Neighbours

```
print("On entire Data")
model_apply(neighbours, X_train, y_train, X_val, y_val, 'KNeighbors')
```


On entire Data

Accuracy: 0.7688645560580815

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.94	0.87	6745
1	0.19	0.05	0.08	1657
accuracy			0.77	8402
macro avg	0.50	0.50	0.48	8402
weighted avg	0.68	0.77	0.71	8402

Confusion Matrix:



```
print("On Random Under Sample")
model_apply(neighbours, X_train_under, y_train_under, X_val_under, y_val_under, 'KNeighbor
```


On Random Under Sample

Accuracy: 0.4939649969824985

Classification Report:

	precision	recall	f1-score	support
0	0.50	0.49	0.49	1677
1	0.49	0.50	0.49	1637
accuracy			0.49	3314
macro avg	0.49	0.49	0.49	3314
weighted avg	0.49	0.49	0.49	3314

Confusion Matrix:



```
print("On SMOTE Sample")
model_apply(neighbours, X_train_over, y_train_over, X_val_over, y_val_over, 'KNeighborsOve
```

On SMOTE Sample

Accuracy: 0.6706946400770999

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

## ▼ Neural Networks

```
def train_neural_networks(model, X_training, y_training, X_validation, y_validation, learning_rate):
    Adam(learning_rate=learning_rate)
    model.compile(
        loss='binary_crossentropy',
        optimizer='adam',
        metrics=[metrics.Recall(), metrics.Precision()]
    )
    return model.fit(X_training, y_training, validation_data=(X_validation, y_validation), epochs=100)
```

```
def show_and_save_model(model, name, X_validation, y_validation):
    y_pred = model.predict(X_validation) > 0.5

    print("Accuracy: {}".format(accuracy_score(y_validation, y_pred)))

    print("Classification Report: \n{}".format(classification_report(y_validation, y_pred)))

    conf_matrix = confusion_matrix(y_validation, y_pred)
    print("Confusion Matrix: ")
    sns.heatmap(conf_matrix, annot=True, fmt='d')

    filename = "model/{}.pkl".format(name)
    pickle.dump(model, open(filename, 'wb'))
```

```
neg_under, pos_under = np.bincount(y_train_under)
```

```
neg_over, pos_over = np.bincount(y_train_over)
```

```
initial_bias_under = np.log([pos_under/neg_under])
initial_bias_under
```

```
array([0.00603502])
```

```
initial_bias_over = np.log([pos_over/neg_over])
initial_bias_over
```

```
array([0.00122324])
```

## ▼ On Random Under Sample

```
model = Sequential()
```

```

model.add(Dense(128, input_shape=(11,), activation='relu'))
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(64, activation='relu'))
model.add(Dense(1, activation='sigmoid', bias_initializer=Constant(initial_bias_under)))

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	1536
dense_1 (Dense)	(None, 128)	16512
dropout (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 1)	65
Total params: 26,369		
Trainable params: 26,369		
Non-trainable params: 0		

```
hist_us = train_neural_networks(model, X_train_under, y_train_under, X_val_under, y_val_ur
```

```

Epoch 1/25
884/884 [=====] - 5s 4ms/step - loss: 0.6969 - recall: 0.52
Epoch 2/25
884/884 [=====] - 3s 4ms/step - loss: 0.6945 - recall: 0.51
Epoch 3/25
884/884 [=====] - 3s 3ms/step - loss: 0.6933 - recall: 0.58
Epoch 4/25
884/884 [=====] - 3s 4ms/step - loss: 0.6931 - recall: 0.61
Epoch 5/25
884/884 [=====] - 3s 4ms/step - loss: 0.6929 - recall: 0.56
Epoch 6/25
884/884 [=====] - 3s 3ms/step - loss: 0.6918 - recall: 0.63
Epoch 7/25
884/884 [=====] - 3s 3ms/step - loss: 0.6925 - recall: 0.63
Epoch 8/25
884/884 [=====] - 3s 4ms/step - loss: 0.6920 - recall: 0.60
Epoch 9/25
884/884 [=====] - 3s 3ms/step - loss: 0.6926 - recall: 0.58
Epoch 10/25
884/884 [=====] - 3s 4ms/step - loss: 0.6919 - recall: 0.53
Epoch 11/25
884/884 [=====] - 3s 4ms/step - loss: 0.6911 - recall: 0.57
Epoch 12/25
884/884 [=====] - 3s 4ms/step - loss: 0.6909 - recall: 0.58
Epoch 13/25
884/884 [=====] - 3s 3ms/step - loss: 0.6908 - recall: 0.60
Epoch 14/25
884/884 [=====] - 3s 3ms/step - loss: 0.6905 - recall: 0.54
Epoch 15/25

```



```

884/884 [=====] - 3s 3ms/step - loss: 0.6898 - recall: 0.60
Epoch 16/25
884/884 [=====] - 3s 3ms/step - loss: 0.6897 - recall: 0.55
Epoch 17/25
884/884 [=====] - 5s 6ms/step - loss: 0.6889 - recall: 0.54
Epoch 18/25
884/884 [=====] - 3s 4ms/step - loss: 0.6883 - recall: 0.57
Epoch 19/25
884/884 [=====] - 3s 4ms/step - loss: 0.6871 - recall: 0.54
Epoch 20/25
884/884 [=====] - 3s 3ms/step - loss: 0.6875 - recall: 0.55
Epoch 21/25
884/884 [=====] - 3s 3ms/step - loss: 0.6859 - recall: 0.52
Epoch 22/25
884/884 [=====] - 4s 4ms/step - loss: 0.6857 - recall: 0.52
Epoch 23/25
884/884 [=====] - 3s 3ms/step - loss: 0.6853 - recall: 0.51
Epoch 24/25
884/884 [=====] - 5s 6ms/step - loss: 0.6835 - recall: 0.54
Epoch 25/25
884/884 [=====] - 6s 7ms/step - loss: 0.6823 - recall: 0.52

```

```
show_and_save_model(model, 'TwoLayerBiasUnder', X_val_under, y_val_under)
```


```
104/104 [=====] - 0s 2ms/step
```

```
Accuracy: 0.48974049487024746
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.50	0.55	0.52	1677
1	0.48	0.43	0.46	1637
accuracy			0.49	3314
macro avg	0.49	0.49	0.49	3314
weighted avg	0.49	0.49	0.49	3314

```
Confusion Matrix:
```



## ▼ On SMOTE Sample

```
model = Sequential()
```

```

model.add(Dense(256, input_shape=(11,), activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(128, activation='relu'))
model.add(Dense(128, activation='relu'))
model.add(Dense(1, activation='sigmoid', bias_initializer=Constant(initial_bias_over)))

model.summary()

```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
=====		
dense_22 (Dense)	(None, 256)	3072
dropout_8 (Dropout)	(None, 256)	0
dense_23 (Dense)	(None, 256)	65792
dropout_9 (Dropout)	(None, 256)	0
dense_24 (Dense)	(None, 128)	32896
dense_25 (Dense)	(None, 128)	16512
dense_26 (Dense)	(None, 1)	129
=====		
Total params: 118,401		
Trainable params: 118,401		
Non-trainable params: 0		
=====		

```

neural_networks(model, X_train_over, y_train_over, X_val_over, y_val_over, learning_rate=0.

```

```

=====] - 7s 7ms/step - loss: 0.5933 - recall_6: 0.6315 - precision_6: 0.7094
=====] - 7s 6ms/step - loss: 0.5935 - recall_6: 0.6389 - precision_6: 0.7035
=====] - 7s 7ms/step - loss: 0.5935 - recall_6: 0.6314 - precision_6: 0.7081
=====] - 7s 7ms/step - loss: 0.5930 - recall_6: 0.6304 - precision_6: 0.7115
=====] - 7s 7ms/step - loss: 0.5929 - recall_6: 0.6351 - precision_6: 0.7097
=====] - 7s 7ms/step - loss: 0.5926 - recall_6: 0.6361 - precision_6: 0.7080
=====] - 8s 7ms/step - loss: 0.5917 - recall_6: 0.6473 - precision_6: 0.7034
=====] - 9s 8ms/step - loss: 0.5911 - recall_6: 0.6463 - precision_6: 0.7048
=====] - 8s 8ms/step - loss: 0.5902 - recall_6: 0.6436 - precision_6: 0.7096
=====] - 7s 7ms/step - loss: 0.5908 - recall_6: 0.6450 - precision_6: 0.7062
=====] - 7s 7ms/step - loss: 0.5906 - recall_6: 0.6461 - precision_6: 0.7095

```

```

=====] - 7s 7ms/step - loss: 0.5901 - recall_6: 0.6461 - precision_6: 0.7047
=====] - 8s 7ms/step - loss: 0.5887 - recall_6: 0.6508 - precision_6: 0.7057
=====] - 7s 7ms/step - loss: 0.5885 - recall_6: 0.6478 - precision_6: 0.7084
=====] - 7s 7ms/step - loss: 0.5898 - recall_6: 0.6457 - precision_6: 0.7074
=====] - 7s 7ms/step - loss: 0.5885 - recall_6: 0.6398 - precision_6: 0.7127
=====] - 7s 7ms/step - loss: 0.5879 - recall_6: 0.6472 - precision_6: 0.7123
=====] - 7s 7ms/step - loss: 0.5881 - recall_6: 0.6561 - precision_6: 0.7066
=====] - 7s 7ms/step - loss: 0.5883 - recall_6: 0.6398 - precision_6: 0.7085
=====] - 7s 7ms/step - loss: 0.5870 - recall_6: 0.6399 - precision_6: 0.7117
=====] - 8s 7ms/step - loss: 0.5868 - recall_6: 0.6420 - precision_6: 0.7134
=====] - 7s 7ms/step - loss: 0.5863 - recall_6: 0.6365 - precision_6: 0.7157
=====] - 7s 7ms/step - loss: 0.5874 - recall_6: 0.6350 - precision_6: 0.7127
=====] - 7s 7ms/step - loss: 0.5862 - recall_6: 0.6456 - precision_6: 0.7118
=====] - 7s 7ms/step - loss: 0.5866 - recall_6: 0.6414 - precision_6: 0.7128
=====] - 7s 7ms/step - loss: 0.5873 - recall_6: 0.6453 - precision_6: 0.7115
=====] - 7s 7ms/step - loss: 0.5865 - recall_6: 0.6483 - precision_6: 0.7123
=====] - 7s 7ms/step - loss: 0.5860 - recall_6: 0.6488 - precision_6: 0.7115
=====] - 7s 7ms/step - loss: 0.5847 - recall_6: 0.6466 - precision_6: 0.7140

```

```
hist_os = train_neural_networks(model, X_train_over, y_train_over, X_val_over, y_val_over,
```

```

loss: 0.5805 - recall_7: 0.6507 - precision_7: 0.7181 - val_loss: 0.5858 - val_reca
loss: 0.5813 - recall_7: 0.6527 - precision_7: 0.7160 - val_loss: 0.5833 - val_reca
loss: 0.5808 - recall_7: 0.6438 - precision_7: 0.7176 - val_loss: 0.5845 - val_reca
loss: 0.5820 - recall_7: 0.6412 - precision_7: 0.7196 - val_loss: 0.5856 - val_reca
loss: 0.5817 - recall_7: 0.6479 - precision_7: 0.7172 - val_loss: 0.5827 - val_reca
loss: 0.5804 - recall_7: 0.6523 - precision_7: 0.7159 - val_loss: 0.5831 - val_reca
loss: 0.5801 - recall_7: 0.6536 - precision_7: 0.7156 - val_loss: 0.5836 - val_reca
loss: 0.5804 - recall_7: 0.6540 - precision_7: 0.7149 - val_loss: 0.5822 - val_reca
loss: 0.5800 - recall_7: 0.6500 - precision_7: 0.7187 - val_loss: 0.5836 - val_reca
loss: 0.5792 - recall_7: 0.6614 - precision_7: 0.7148 - val_loss: 0.5829 - val_reca

```

```

loss: 0.5707 - recall_7: 0.6542 - precision_7: 0.7174 - val_loss: 0.5842 - val_reca

```

```
loss: 0.5797 - recall_7: 0.6542 - precision_7: 0.7174 - val_loss: 0.5843 - val_reca
loss: 0.5800 - recall_7: 0.6542 - precision_7: 0.7168 - val_loss: 0.5817 - val_reca
loss: 0.5809 - recall_7: 0.6507 - precision_7: 0.7160 - val_loss: 0.5857 - val_reca
loss: 0.5800 - recall_7: 0.6597 - precision_7: 0.7143 - val_loss: 0.5824 - val_reca
loss: 0.5796 - recall_7: 0.6490 - precision_7: 0.7202 - val_loss: 0.5837 - val_reca
loss: 0.5802 - recall_7: 0.6487 - precision_7: 0.7207 - val_loss: 0.5858 - val_reca
loss: 0.5786 - recall_7: 0.6545 - precision_7: 0.7166 - val_loss: 0.5827 - val_reca
loss: 0.5795 - recall_7: 0.6471 - precision_7: 0.7203 - val_loss: 0.5855 - val_reca
loss: 0.5798 - recall_7: 0.6460 - precision_7: 0.7219 - val_loss: 0.5826 - val_reca
loss: 0.5791 - recall_7: 0.6481 - precision_7: 0.7233 - val_loss: 0.5824 - val_reca
loss: 0.5793 - recall_7: 0.6536 - precision_7: 0.7176 - val_loss: 0.5826 - val_reca
loss: 0.5779 - recall_7: 0.6517 - precision_7: 0.7200 - val_loss: 0.5826 - val_reca
loss: 0.5778 - recall_7: 0.6591 - precision_7: 0.7187 - val_loss: 0.5830 - val_reca
loss: 0.5779 - recall_7: 0.6544 - precision_7: 0.7175 - val_loss: 0.5832 - val_reca
loss: 0.5800 - recall_7: 0.6506 - precision_7: 0.7172 - val_loss: 0.5818 - val_reca
loss: 0.5779 - recall_7: 0.6526 - precision_7: 0.7209 - val_loss: 0.5816 - val_reca
loss: 0.5787 - recall_7: 0.6581 - precision_7: 0.7180 - val_loss: 0.5854 - val_reca
loss: 0.5777 - recall_7: 0.6601 - precision_7: 0.7139 - val_loss: 0.5836 - val_reca
loss: 0.5784 - recall_7: 0.6584 - precision_7: 0.7138 - val_loss: 0.5821 - val_reca
```

```
show_and_save_model(model, 'TwoLayerBiasOver', X_val_over, y_val_over)
```

422/422 [=====] - 1s 2ms/step

Accuracy: 0.6959003632589518

Classification Report:

	precision	recall	f1-score	support
0	0.71	0.67	0.69	6761
1	0.69	0.72	0.70	6728
accuracy			0.70	13489

## ▼ Prediction on Test Model

```
test_data.drop('ID', axis=1, inplace=True)
```

```
categorical_cols = [col for col in test_data.columns if test_data[col].dtype=="O"]
```

```
label_encoder = LabelEncoder()
```

```
for col in categorical_cols:
    test_data[col] = label_encoder.fit_transform(test_data[col])
```

```
test_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28000 entries, 0 to 27999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Income                28000 non-null  int64
1   Age                   28000 non-null  int64
2   Experience             28000 non-null  int64
3   Married/Single        28000 non-null  int64
4   House_Ownership       28000 non-null  int64
5   Car_Ownership         28000 non-null  int64
6   Profession            28000 non-null  int64
7   CITY                  28000 non-null  int64
8   STATE                 28000 non-null  int64
9   CURRENT_JOB_YRS       28000 non-null  int64
10  CURRENT_HOUSE_YRS     28000 non-null  int64
dtypes: int64(11)
memory usage: 2.3 MB
```

```
cols = test_data.columns
```

```
test_data = scaler.transform(test_data.to_numpy())
test_data = pd.DataFrame(test_data, columns=list(cols))
test_data
```

	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownersh
<b>0</b>	0.837459	0.531939	1.506277	0.332830	0.283607	-0.6585
<b>1</b>	-1.312342	-1.460938	-0.811160	0.332830	0.283607	-0.6585
<b>2</b>	1.362288	0.004413	0.347558	0.332830	0.283607	-0.6585
<b>3</b>	-1.058526	-0.054201	-0.149036	-3.004535	0.283607	1.5184
<b>4</b>	-1.730457	-1.460938	1.340745	0.332830	0.283607	1.5184
...	...	...	...	...	...	...
<b>27995</b>	1.729099	0.414711	0.513089	0.332830	0.283607	-0.6585
<b>27996</b>	-0.719829	-0.171429	-0.149036	0.332830	0.283607	-0.6585
<b>27997</b>	1.077325	-1.519552	-0.811160	0.332830	0.283607	-0.6585
<b>27998</b>	1.561620	0.063027	0.513089	0.332830	0.283607	1.5184
<b>27999</b>	1.483733	-0.464499	-0.149036	0.332830	0.283607	-0.6585

28000 rows × 11 columns

```
def load_and_test(filename):
    model = pickle.load(open(filename, 'rb'))
    y_pred = model.predict(test_data.iloc[:, : 11])
    return y_pred
```

## ▼ Logistic Regression

```
test_data['lr_predict'] = list(load_and_test('model/LogisticRegressionOver.pkl'))
test_data
```

	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownersh
0	0.837459	0.531939	1.506277	0.332830	0.283607	-0.6585
1	-1.312342	-1.460938	-0.811160	0.332830	0.283607	-0.6585

```
test_data['lr_predict'].value_counts()
```

```
1    14850
0    13150
Name: lr_predict, dtype: int64
```

▼ Decision Tree

27996	-0.719829	-0.171429	-0.149036	0.332830	0.283607	-0.6585
-------	-----------	-----------	-----------	----------	----------	---------

```
test_data['dt_predict'] = list(load_and_test('model/DecisionTreeOver.pkl'))
test_data
```

	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownersh
0	0.837459	0.531939	1.506277	0.332830	0.283607	-0.6585
1	-1.312342	-1.460938	-0.811160	0.332830	0.283607	-0.6585
2	1.362288	0.004413	0.347558	0.332830	0.283607	-0.6585
3	-1.058526	-0.054201	-0.149036	-3.004535	0.283607	1.5184
4	-1.730457	-1.460938	1.340745	0.332830	0.283607	1.5184
...	...	...	...	...	...	...
27995	1.729099	0.414711	0.513089	0.332830	0.283607	-0.6585
27996	-0.719829	-0.171429	-0.149036	0.332830	0.283607	-0.6585
27997	1.077325	-1.519552	-0.811160	0.332830	0.283607	-0.6585
27998	1.561620	0.063027	0.513089	0.332830	0.283607	1.5184
27999	1.483733	-0.464499	-0.149036	0.332830	0.283607	-0.6585

28000 rows × 13 columns



--

```
test_data['dt_predict'].value_counts()
```

```
1    16859
0    11141
Name: dt_predict, dtype: int64
```


▼ Random Forest

```
test_data['rf_predict'] = list(load_and_test('model/RandomForestOver.pkl'))
```

```
test_data['rf_predict'] = list(load_and_test('model/RandomForestOver.pkl'))
test_data
```

	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownersh
<b>0</b>	0.837459	0.531939	1.506277	0.332830	0.283607	-0.6585
<b>1</b>	-1.312342	-1.460938	-0.811160	0.332830	0.283607	-0.6585
<b>2</b>	1.362288	0.004413	0.347558	0.332830	0.283607	-0.6585
<b>3</b>	-1.058526	-0.054201	-0.149036	-3.004535	0.283607	1.5184
<b>4</b>	-1.730457	-1.460938	1.340745	0.332830	0.283607	1.5184
...	...	...	...	...	...	...
<b>27995</b>	1.729099	0.414711	0.513089	0.332830	0.283607	-0.6585
<b>27996</b>	-0.719829	-0.171429	-0.149036	0.332830	0.283607	-0.6585
<b>27997</b>	1.077325	-1.519552	-0.811160	0.332830	0.283607	-0.6585
<b>27998</b>	1.561620	0.063027	0.513089	0.332830	0.283607	1.5184
<b>27999</b>	1.483733	-0.464499	-0.149036	0.332830	0.283607	-0.6585

28000 rows × 14 columns



```
test_data['rf_predict'].value_counts()
```

```
0    15453
1    12547
Name: rf_predict, dtype: int64
```

## ▼ K Nearest Neighbours

```
test_data['knn_predict'] = list(load_and_test('model/KNeighborsOver.pkl'))
test_data
```



	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownersh
0	0.837459	0.531939	1.506277	0.332830	0.283607	-0.6585
1	-1.312342	-1.460938	-0.811160	0.332830	0.283607	-0.6585
2	1.362288	0.004413	0.347558	0.332830	0.283607	-0.6585
3	-1.058526	-0.054201	-0.149036	-3.004535	0.283607	1.5184
4	-1.730457	-1.460938	1.340745	0.332830	0.283607	1.5184
...	...	...	...	...	...	...
27995	1.729099	0.414711	0.513089	0.332830	0.283607	-0.6585

```
test_data['knn_predict'].value_counts()
```

```
0    15955
```

```
1    12045
```

```
Name: knn_predict, dtype: int64
```

```
27996 1.492722 -0.464499 -0.149036 0.332830 0.283607 0.6585
```

## ▼ Two Layer with Bias Network

```
test_data['b_predict'] = [1 if item else 0 for sublist in list(load_and_test('model/TwoLayer', test_data
```

```
875/875 [=====] - 2s 2ms/step
```

	Income	Age	Experience	Married/Single	House_Ownership	Car_Ownersh
0	0.837459	0.531939	1.506277	0.332830	0.283607	-0.6585
1	-1.312342	-1.460938	-0.811160	0.332830	0.283607	-0.6585
2	1.362288	0.004413	0.347558	0.332830	0.283607	-0.6585
3	-1.058526	-0.054201	-0.149036	-3.004535	0.283607	1.5184
4	-1.730457	-1.460938	1.340745	0.332830	0.283607	1.5184
...	...	...	...	...	...	...
27995	1.729099	0.414711	0.513089	0.332830	0.283607	-0.6585
27996	-0.719829	-0.171429	-0.149036	0.332830	0.283607	-0.6585
27997	1.077325	-1.519552	-0.811160	0.332830	0.283607	-0.6585
27998	1.561620	0.063027	0.513089	0.332830	0.283607	1.5184
27999	1.483733	-0.464499	-0.149036	0.332830	0.283607	-0.6585

```
28000 rows × 16 columns
```


```
test_data['b_predict'].value_counts()
```

```
1    14700
0    13300
Name: b_predict, dtype: int64
```

## ▼ Saving Output Dataset

```
predicted_data = pd.DataFrame({
    'Id': list(test_data.index + 1)
})
```

```
predicted_data = pd.concat([predicted_data, test_data.iloc[:, 11:]], axis=1)
predicted_data
```

	Id	lr_predict	dt_predict	rf_predict	knn_predict	b_predict	
<b>0</b>	1	0	1	1	0	1	
<b>1</b>	2	1	1	1	1	1	
<b>2</b>	3	0	1	0	0	1	
<b>3</b>	4	0	0	0	1	1	
<b>4</b>	5	0	1	0	0	0	
...	...	...	...	...	...	...	
<b>27995</b>	27996	1	1	1	1	1	
<b>27996</b>	27997	0	1	0	0	0	
<b>27997</b>	27998	1	0	0	1	1	
<b>27998</b>	27999	0	1	0	0	0	
<b>27999</b>	28000	1	1	1	1	1	

28000 rows × 6 columns

```
predicted_data.to_csv('data/Prediction on Test.csv')
```

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 1:22 PM

